

## 1.Importing the headers

```
In [1]: import io
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import statistics as st
import sklearn
```

## 2.Downloading and loading the churn\_modelling dataset

```
In [2]: from google.colab import files
uploaded = files.upload()
```

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.  
Saving Churn\_Modelling.csv to Churn\_Modelling.csv

```
Out [3]: dataset = pd.read_csv(io.BytesIO(uploaded['Churn_Modelling.csv']))
```

```
In [72]: df = pd.read_csv('Churn_Modelling.csv')
df.head()
```

```
Out [72]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1	101348.08	1
1	2	15647311	Hair	608	Spain	Female	41	1	83807.86	1	0	1	112542.56	0
2	3	15616409	Ouin	502	France	Female	42	8	159660.80	3	1	0	113931.57	1
3	4	15701324	Burn	699	France	Female	39	1	0.00	2	0	0	92626.63	0
4	5	15757688	Michel	850	Spain	Female	43	2	125510.82	1	1	1	79064.10	0

```
In [73]: df = df.drop(columns=['RowNumber', 'CustomerId', 'Surname'])
df.head()
```

```
Out [73]:
```


	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	619	France	Female	42	2	0.00	1	1	1	101348.08	1
1	608	Spain	Female	41	1	83807.86	1	0	1	112542.56	0
2	502	France	Female	42	8	159660.80	3	1	0	113931.57	1
3	699	France	Female	39	1	0.00	2	0	0	92626.63	0
4	850	Spain	Female	43	2	125510.82	1	1	1	79064.10	0

2	502	France	Female	42	0	109600.80	2	1	0	112831.37	1
3	489	France	Female	39	1	0.00	2	0	0	81825.63	0
4	850	Spain	Female	43	2	125510.82	1	1	1	79084.10	0

```
In [74]: df['IsActiveMember'] = df['IsActiveMember'].astype('category')
df['Exited'] = df['Exited'].astype('category')
df['HasCrCard'] = df['HasCrCard'].astype('category')
```

### 3a.Univariate Analysis

```
In [75]: sns.kdeplot(x='CreditScore', data = df, hue = 'Exited')
plt.show()
```



### 3b.Bivariate Analysis

```
In [76]: density = df['Exited'].value_counts(normalize=True).reset_index()
sns.barplot(data=density, x='index', y='Exited', )
density
```

```
Out[76]:
```

	index	Exited
0	0	0.7961
1	1	0.2037




### 3b. Bivariate Analysis

```
In [76]: density = df['Exited'].value_counts(normalize=True).reset_index()
sns.barplot(data=density, x='index', y='Exited', )
density
```

```
Out[76]:
```

	index	Exited
0	0	0.7963
1	1	0.2037

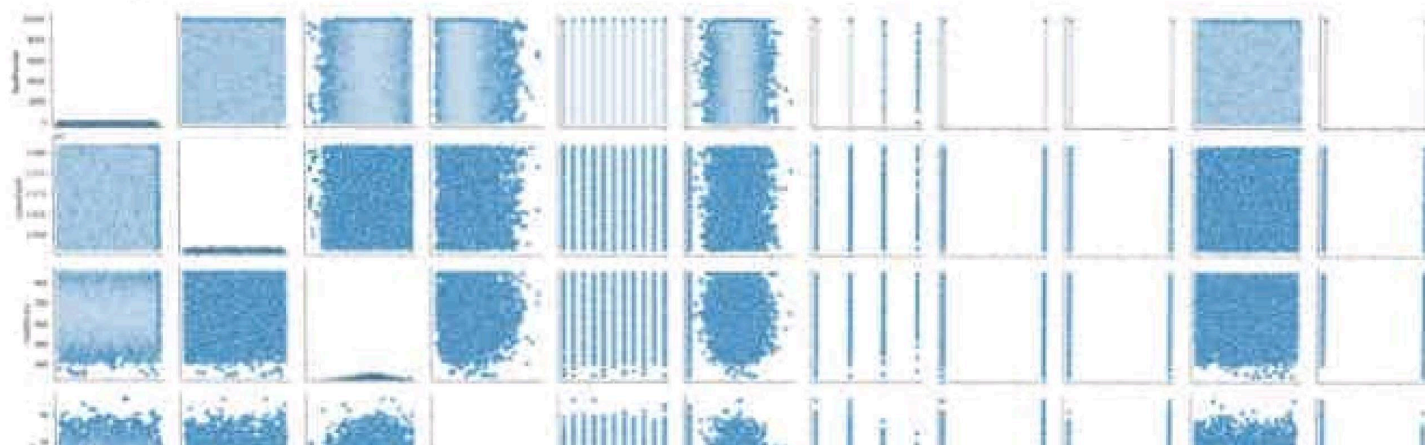


### 3c. Multivariate Analysis

```
In [77]: sm.pairplot(dataset)
```

```
Out[77]:
```

seaborn.axisgrid.PairGrid at 0c7f3908c57f90>





## 4.Descriptive Statistics **bold text**

In [78]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 11 columns):
 #   Column             Non-Null Count  Dtype  
---  --
 0   CreditScore         10000 non-null  int64  
 1   Geography           10000 non-null  object  
 2   Gender              10000 non-null  object  
 3   Age                 10000 non-null  int64  
 4   Tenure              10000 non-null  int64  
 5   Balance             10000 non-null  float64  
 6   NumOfProducts       10000 non-null  int64  
 7   HasCrCard           10000 non-null  category
 8   IsActiveMember      10000 non-null  category
 9   EstimatedSalary     10000 non-null  float64  
10   Exited              10000 non-null  category
dtypes: category(3), float64(2), int64(4), object(2)
memory usage: 654.8+ KB
```

In [79]:

```
df.describe()
```

Out[79]:

	CreditScore	Age	Tenure	Balance	NumOfProducts	EstimatedSalary
count	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000
mean	652.528000	38.821800	5.012800	76485.889208	1.530200	90096.236881
std	90.651299	11.487906	2.852174	62387.405202	0.581654	17110.482218
min	350.000000	18.000000	0.000000	0.000000	1.000000	11.580000
25%	584.000000	32.000000	1.000000	0.000000	1.000000	51002.110000
50%	652.000000	37.000000	5.000000	97198.540000	1.000000	902182.915000
75%	718.000000	44.000000	7.000000	127644.240000	2.000000	148388.247500
max	850.000000	62.000000	10.000000	250986.090000	4.000000	199492.480000

## 5.Handle Missing Values

In [80]:

```
df.isna().sum()
```

Out[80]:

```
CreditScore      0
Geography         0
Gender           0
Age              0
Tenure           0
Balance          0
NumOfProducts    0
HasCrCard        0
IsActiveMember   0
EstimatedSalary  0
```




## 6. Find the Outliers and Replace the Outlier

### Finding outliers

```
In [84]: def box_scatter(data, x, y):  
    fig, (ax1, ax2) = plt.subplots(nrows=2, ncols=1, figsize=(10,6))  
    sns.boxplot(data=data, x=x, ax=ax1)  
    sns.scatterplot(data=data, x=x, y=y, ax=ax2)
```

```
In [85]: box_scatter(df, 'CreditScore', 'Exited');  
plt.tight_layout()  
print(f"# of Riverista Outliers: {len(df.loc[df['CreditScore'] < 400])}")
```


# of Riverista Outliers: 19



```
In [86]: box_scatter(df, 'Age', 'Exited');  
plt.tight_layout()  
print(f"# of Riverista Outliers: {len(df.loc[df['Age'] > 87])}")
```



# of Riverista Outliers: 1







```
In [34]: box_scatter(df, 'Balance', 'Exited');
plt.tight_layout()
print(f"% of Bivariate Outliers: {len(df.loc[df['Balance'] > 220000])}")
```

% of Bivariate Outliers: 4



```
In [47]: box_scatter(df, 'EstimatedSalary', 'Exited');
plt.tight_layout()
```



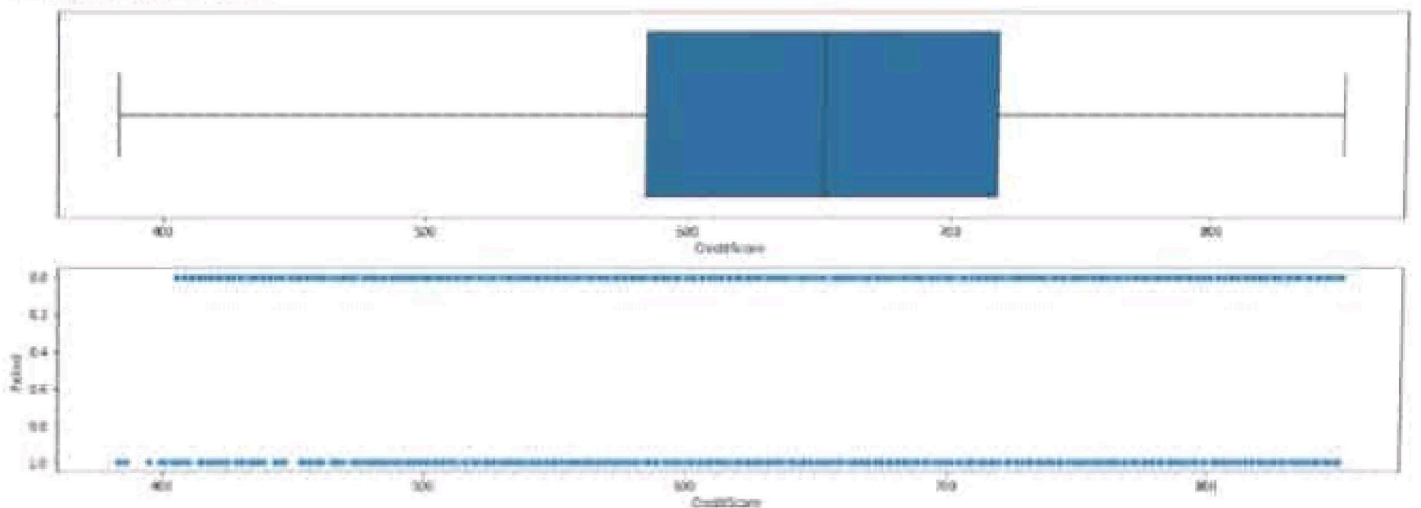


## Removing the Outliers

```
In [86]: for i in df:
         if df[i].dtype=='int64' or df[i].dtype=='float64':
             q1=df[i].quantile(0.25)
             q3=df[i].quantile(0.75)
             iqr=q3-q1
             upper=q3+1.5*iqr
             lower=q1-1.5*iqr
             df[i]=df[i].where(df[i] > upper, upper, df[i])
             df[i]=df[i].where(df[i] < lower, lower, df[i])

In [87]: box_scatter(df, 'CreditScore', 'Exited')
plt.tight_layout()
print(f"# of Bivariate Outliers: {len(df.loc[df['CreditScore'] < 400])}")
```

# of Bivariate Outliers: 15






400 500 600 700 800  
DebtScore


```
In [87]: box_scatter(df, 'Age', 'Exited');  
plt.tight_layout()  
print(f"% of Bivariate Outliers: {len(df.loc[df['Age'] > 87])}")
```

% of Bivariate Outliers: 8



```
In [88]: box_scatter(df, 'Balance', 'Exited');  
plt.tight_layout()  
print(f"% of Bivariate Outliers: {len(df.loc[df['Balance'] > 22000])}")
```

% of Bivariate Outliers: 4



## 7.Categorical Encoding

```
In [14]: from sklearn.preprocessing import LabelEncoder
encoder=LabelEncoder()
for i in df:
    if df[i].dtype=="object":
        df[i]=encoder.fit_transform(df[i])
```

## 8.Split the data into dependent and independent variables

```
In [15]: X = dataset
X.drop(['RowNumber', 'Exited'], axis=1)

Y = pd.DataFrame(dataset['Exited'])
```

```
In [16]: print (X.columns)
print (Y.columns)

Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore', 'Geography',
       'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',
       'IsActiveMember', 'EstimatedSalary', 'Exited'],
      dtype='object')
Index(['Exited'], dtype='object')
```

## 9.Scale and independent variables

```
In [18]: print (Y)

      Exited
0         1
1         0
2         1
3         0
4         0
...
9995      0
9996      0
9997      1
9998      1
9999      0

[10000 rows x 1 columns]
```

```
In [19]: from sklearn import preprocessing
normalized_Y = preprocessing.normalize(Y)
print (normalized_Y)
```

```
[[1.]
 [0.]
 [1.]
 ...
 [1.]
 [1.]
 [0.]]
```

In [18]:

```
standard_Y = Y.copy()

from sklearn import preprocessing

ss = preprocessing.StandardScaler()
ss.fit(standard_Y)

print (standard_Y)
```

```
      Exited
0         1
1         0
2         1
3         0
4         0
...
9995      0
9996      0
9997      1
9998      1
9999      0
```

[10000 rows x 1 columns]

## 10.Perform Train Test Split

In [19]:

```
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(X, Y)
```

In [20]:

```
x_train
```

Out[20]:

RowIndex	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited	
6811	4812	15771222	Cigujeher	779	France	Female	42	5	0.00	2	0	0	25951.91	0
2974	2975	15746731	Eamex	583	Germany	Male	58	8	88128.58	1	1	1	170705.53	0
952	953	11800149	Wuxul	571	Germany	Female	46	8	111577.01	1	0	1	188271.90	0
5643	5644	15782096	Vullone	616	Spain	Female	36	6	0.00	4	1	1	12816.32	1
7061	7062	15789201	Thomson	603	Germany	Female	35	9	545623.36	1	1	0	513181.62	0
...	...	...	...	...	...	...	...	...	...	...	...	...	...	
4653	4654	15579017	Sinclair	488	France	Female	51	3	0.00	2	0	1	174086.28	1
1317	1318	15720702	Shih	789	France	Male	37	3	0.00	1	1	0	121883.87	1

In [106]:

x\_test

Out[106]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
7220	7221	15706637	Cheng	718	Spain	Male	40	9	0.00	2	0	0	121537.81	
1154	1155	15706887	Iedanzu	538	Spain	Male	39	2	122773.30	2	1	1	58467.08	
3156	3157	15509641	Sung	692	Germany	Female	41	6	110701.29	1	1	0	59354.24	
4514	4515	15658670	Chim	669	France	Female	38	6	0.00	2	1	0	64048.16	
2817	2818	15708244	Maclean	538	Spain	Female	30	8	0.00	2	1	1	41192.95	
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
3443	3444	15605710	Morero	597	Spain	Female	46	4	0.00	2	1	0	58607.16	
7123	7124	15682686	Chukwemeka	722	France	Female	38	3	0.00	2	0	1	167984.72	
1997	1998	15624781	Mbenke	672	France	Female	34	1	142151.25	2	1	1	163753.24	
5735	5736	15756079	Greenwood	585	Spain	Female	44	6	0.00	2	0	1	101728.46	
4348	4349	15574387	Ni	625	Germany	Female	44	2	79064.85	2	0	1	113291.75	

1500 rows x 14 columns

4

In [107]:

y\_train

Out[107]:

	Exited
6811	0
2874	0
952	0
5643	1
7061	0
...	...
4633	1
1317	1
3661	0
1002	0
5961	0

7500 rows x 1 columns

In [108]:

y\_test

```
952      0
5643     1
7061     0
--
4653     1
1317     1
3661     0
1002     0
5961     0
```

7500 rows x 1 columns

In [19]:

```
y_test
```

Out[19]:

```
Exited
7220     0
1154     0
3154     1
4514     0
2617     0
--
3443     1
7523     0
1997     0
5725     0
4348     0
```

2500 rows x 1 columns

In [ ]: